

Automatic Image Analysis Script

May 6, 2026

Contents

1	Introduction to Automatic Image Analysis	2
1.1	Lecture Aim and the Inverse Problem	2
1.2	From Pixels to Interpretation	2
1.3	Models, Data, and Evaluation	4
1.4	Course Roadmap	5
2	Core Topics	5
2.1	Image Formation and Representation	6
2.2	Learning-Based Methods	6
2.3	Visual Object Models	6
2.4	Probabilistic Estimation	6
3	Exercises	6

1 Introduction to Automatic Image Analysis

The first lecture should give students a mental model for the whole course. We start from the central difficulty of vision, then move through the main processing levels, the recurring goals of the field, and the practical question of how to evaluate whether a vision system is actually reliable.

1.1 Lecture Aim and the Inverse Problem

This introductory lecture frames automatic image analysis as the problem of turning measurements into explanations. A digital image is easy to store: it is a grid of values. It is much harder to understand: we want to infer objects, surfaces, motion, geometry, lighting, and meaning from those values.

The central message for the first session is therefore simple: computer vision is not just image enhancement. It is inference under uncertainty. The image is evidence; the scene that produced it is hidden.




Figure 1: Computer vision tries to invert the image formation process: pixels are observed, but the scene explanation is hidden.

Human vision feels immediate, but this ease is misleading. A camera records a two-dimensional projection of a three-dimensional world. During that projection, many causes are mixed together: object shape, material, illumination, viewpoint, occlusion, sensor noise, and motion. Different scenes can produce very similar images, so there is usually no unique answer that can be read directly from the pixels.

This makes vision an inverse problem. The forward direction is comparatively clear: if we know the scene, the light, the camera, and the materials, we can model how an image is formed. The inverse direction asks us to go backward from image measurements to a plausible explanation of the world. That backward step requires assumptions, models, data, and algorithms.

A compact way to write the forward model is

$$I(x, y) = f(\Pi(X, Y, Z), L, M, C) + \varepsilon, \quad (1)$$

where $I(x, y)$ is the measured pixel value, Π is the camera projection from 3D to 2D, L describes illumination, M material and surface properties, C camera response, and ε sensor noise. Vision asks for some hidden scene state S from the image I :

$$\hat{S} = \arg \max_S p(S | I) = \arg \max_S p(I | S) p(S). \quad (2)$$

The likelihood $p(I | S)$ says which images a scene would probably produce. The prior $p(S)$ encodes what kinds of scenes are plausible before seeing the image.

Instructor cue. Show a familiar image and ask: what would the pixel array alone tell us? Then ask what a person immediately infers: objects, depth, material, intent, and context. The gap between those two answers is the subject of the course.

1.2 From Pixels to Interpretation

Automatic image analysis can be organized as a progression from low-level to high-level tasks.




Figure 2: Ambiguity is unavoidable: different scenes can lead to similar image measurements. Priors and additional evidence help choose an explanation.

- Low-level processing improves or transforms image measurements, for example denoising, filtering, contrast changes, and edge extraction.
- Mid-level analysis groups measurements into more meaningful structures, such as regions, contours, features, motion fields, or depth estimates.
- High-level interpretation assigns semantic meaning, such as recognizing objects, estimating scene layout, detecting abnormalities, or tracking actions.

These levels are useful for teaching, but real systems often mix them. A modern recognition model may learn low-level filters and high-level categories inside one network. A classical pipeline may use carefully designed features followed by statistical inference. In both cases, the goal is the same: reduce the ambiguity of image data until a useful decision or description becomes possible.

Level	Typical question	Examples
Low-level	What local measurements are reliable?	denoising, gradients, edges
Mid-level	Which measurements belong together?	regions, contours, motion, depth
High-level	What does the image mean?	objects, scene layout, actions

Table 1: A teaching taxonomy for image-analysis tasks.

One common low-level operation is convolution:

$$(I * k)(x, y) = \sum_{u=-r}^r \sum_{v=-r}^r I(x - u, y - v) k(u, v). \quad (3)$$

The same mathematical form can blur, sharpen, estimate derivatives, or detect local structure depending on the kernel k .

The course can be read through three recurring goals: recognition, reconstruction, and reorganization.

Recognition. Recognition asks what is present and where it is. Examples include classification, detection, segmentation, reading text, identifying faces, or finding defects in industrial inspection.

Reconstruction. Reconstruction asks what spatial or physical structure produced the image. Examples include estimating depth, camera motion, three-dimensional shape, or the geometry of a scene from one or more images.

Reorganization. Reorganization changes the image representation so that later analysis becomes easier or the image becomes more useful to a person. Examples include registration, stitching, stabilization, denoising, super-resolution, and computational photography.

These goals are connected. Recognition benefits from good representations; reconstruction needs reliable correspondences; reorganization often depends on understanding motion, geometry, or image formation.




Figure 3: Recognition, reconstruction, and reorganization are not isolated topics; each goal supplies information that can support the others.

1.3 Models, Data, and Evaluation

Because images are ambiguous, vision systems need prior knowledge. Sometimes that knowledge is explicit: camera geometry, physical image formation, smooth surfaces, rigid motion, or probabilistic noise models. Sometimes it is learned from data: neural networks discover features and decision rules from many examples.

Classical computer vision often emphasizes explicit models and engineered features. Modern deep learning often emphasizes learned representations and large datasets. The practical lesson is not that one replaces the other completely. Robust systems usually combine several kinds of knowledge: scientific models for image formation, statistical models for uncertainty, engineering choices for reliability, and data-driven learning for complex patterns that are hard to specify by hand.

Source of knowledge	What it contributes	Risk if missing
Physical model	Camera geometry, illumination, projection	Results may violate how images are formed
Statistical model	Noise, uncertainty, priors, inference	The system may overtrust ambiguous measurements
Engineering design	Robustness, speed, failure handling	The method may work only in demonstrations
Data-driven learning	Features and decision rules from examples	Complex visual variation may be underspecified

Table 2: A practical vision system usually combines several kinds of knowledge.

Computer vision is now a basic technology in many domains. Mobile phones use it for autofocus, panorama stitching, image stabilization, and computational photography. Medical systems use

image analysis to support diagnosis and planning. Robots and autonomous vehicles use visual input for localization, mapping, and obstacle understanding. Manufacturing systems inspect parts, read codes, and detect defects. Media tools use vision for tracking, editing, visual effects, and three-dimensional capture.

These applications also show why evaluation matters. A result that looks plausible on one example is not enough. Vision algorithms must be tested against noise, changing illumination, clutter, unusual viewpoints, occlusion, and data that differ from the training or development examples.




Figure 4: Evaluation should move from controlled correctness toward real-world robustness.




Figure 5: A compact summary of the introductory framing: measurements become useful outputs only after they are combined with assumptions, models, and data.

1.4 Course Roadmap

The first part of the course builds the foundations: image representation, filtering, features, transformations, and the basic language of image formation. The next part introduces learning-based methods, especially neural networks, as a way to learn visual representations from data. Later topics connect these ideas to object models, probabilistic estimation, and higher-level recognition or reconstruction tasks.

Throughout the course, keep returning to three questions:

1. What information is actually present in the image measurements?
2. What assumptions are needed to infer the hidden scene or decision?
3. How can we test whether the algorithm works beyond the example in front of us?

Automatic image analysis starts with pixels but does not end there. Its real task is to connect images to explanations: what is in the scene, how it is arranged, how it changes, and what action or decision should follow. The rest of the course develops the tools needed to make that connection systematically.

2 Core Topics

The following sections collect the technical building blocks that will be expanded throughout the script. Each topic should connect back to the central question from the introduction: what can be inferred from image measurements, and what assumptions make that inference possible?

2.1 Image Formation and Representation

Add notes on image acquisition, sampling, color spaces, filtering, and basic image transformations.

A pinhole camera model is often the first geometric abstraction:

$$\lambda \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = K \begin{bmatrix} R & t \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix}. \quad (4)$$

It states that a 3D point (X, Y, Z) is projected to image coordinates (x, y) through camera intrinsics K and pose (R, t) , up to an unknown scale λ .

2.2 Learning-Based Methods

Add notes on neural networks, model training, evaluation, and uncertainty.

For supervised recognition, the training set supplies pairs (I_i, y_i) and the model parameters θ are chosen by minimizing an empirical loss:

$$\hat{\theta} = \arg \min_{\theta} \frac{1}{N} \sum_{i=1}^N \ell(g_{\theta}(I_i), y_i). \quad (5)$$

This shifts part of the modeling burden from hand-designed rules to data, architecture, and evaluation.

2.3 Visual Object Models

Add notes on feature representations, object models, detection, and recognition.

2.4 Probabilistic Estimation

Add notes on Bayesian reasoning, probabilistic models, and parameter estimation.

3 Exercises

Use this section for worked exercise notes and derivations.